



Adaptive composition in dynamic service environments



Lina Barakat*, Simon Miles, Michael Luck

Department of Informatics, King's College London, WC2R 2LS, UK

HIGHLIGHTS

- An adaptive service composition behaviour is proposed.
- The adaptation is instantiated at the earliest possibility, *parallel to execution*.
- This early adaptation is triggered for both corrective and optimising purposes.
- Changes encountered during execution go through a novel *prioritisation* process.
- The service re-selection algorithm is *light* and *optimality retaining*.

ARTICLE INFO

Article history:

Received 24 June 2016

Received in revised form

5 November 2016

Accepted 2 December 2016

Available online 19 December 2016

Keywords:

Service composition

Adaptive service execution

Quality of service

Request-based dominance

ABSTRACT

Due to distribution, participant autonomy and lack of local control, service-based systems operate in highly dynamic and uncertain environments. In the face of such dynamism and volatility, the ability to manage service changes and exceptions during composite service execution is a vital requirement. Most current adaptive composition approaches, however, fail to address service changes without causing undesirable disruptions in execution or considerably degrading the quality of the composite application. In response, this paper presents a novel adaptive execution approach, which efficiently handles service changes occurring at execution time, for both repair and optimisation purposes. The adaptation is performed as soon as possible and in parallel with the execution process, thus reducing interruption time, increasing the chance of a successful recovery, and producing the most optimal solution according to the current environment state. The effectiveness of the proposed approach is demonstrated both analytically and empirically through a case study evaluation applied in the framework of learning object composition. In particular, the results show that, even with frequent changes (e.g. 20 changes per service execution), or in the cases where interference with execution is non-preventable (e.g., when an executed service delivers unanticipated quality values), our approach manages to recover from the situation with minimal interruption.

© 2016 Elsevier B.V. All rights reserved.

1. Introduction

Service-oriented computing (SOC) is a suitable paradigm for the sharing of resources and functionalities in large-scale open distributed environments (e.g. the web, computational Grids, and peer-to-peer systems). In this paradigm, providers encapsulate their offerings, ranging from expensive hardware components to entire applications, within *services* and expose them through uniform, machine-readable interfaces (or metadata) on a network of

customers. Via their accessibility, reusability, and loose coupling, services provide the building blocks for rapid and low-cost development of complex distributed applications spanning organisational boundaries. A key feature enabled by SOC is the dynamic binding mechanism. Based on this, a composite application (e.g. a business process, scientific workflow, or e-learning experience) can be structured as a collection of interdependent *abstract* tasks, with *concrete* services being selected for these tasks at run time according to service availability and specific user quality of service (QoS) needs, thus achieving great flexibility and personalisation.

Open distributed service-based systems, however, exhibit high degrees of dynamism and uncertainty for several reasons, either intentional or unintentional. Specifically, existing service providers, being autonomous and self-interested, may choose not to fulfil their promises (e.g. announce false capabilities to attract more customers), to upgrade/degrade their quality

* Corresponding author.

E-mail addresses: lina.barakat@kcl.ac.uk (L. Barakat), simon.miles@kcl.ac.uk (S. Miles), michael.luck@kcl.ac.uk (M. Luck).

<http://dx.doi.org/10.1016/j.future.2016.12.003>

0167-739X/© 2016 Elsevier B.V. All rights reserved.

offerings (e.g. driven by competition), or to disconnect from the system at any time, while new providers could join instead. Even with long-standing and cooperative providers, availability and quality estimates of services could still frequently change due to other factors. For instance, a service's response time could be significantly affected by the provider's load and network traffic at that moment. Similarly, a service might suddenly become unavailable due to network/hardware failure.

Although the dynamic binding of services offers some tolerance against such dynamism and uncertainty, it does not guarantee the successful execution of the composite application, i.e. that the selected component services, for composition, behave as expected. This is because the selection step normally takes place before the start of execution, i.e. services are selected for *all* tasks in advance, to reason effectively about the satisfaction of global (application-level) quality criteria (e.g. total price and total time). Hence, changes to a selected service could occur at any time before the actual invocation of this service, especially when executing complex applications involving many tasks (the case with most realistic applications), or better services could emerge, making the selected service combination no longer valid or no longer the best option.

Consequently, to accommodate service volatility, it is essential for the composite application to be equipped with *adaptation capabilities* at execution time. This is especially important for increasing the satisfaction of the end-user (customer), by improving the quality of the selected solution, and eliminating any inefficiencies from the customer perspective. In light of this, adaptation capabilities should include the following goals: **(G1)** recovering from unexpected situations on their occurrence, so that the application continues its intended execution, or at least terminates in a consistent state, in spite of the occurrence of failure or violation; **(G2)** exploiting new emerging opportunities to enhance the quality of the selected solution at any execution stage; **(G3)** proactively preventing future breaches and faulty behaviour by performing early corrective actions, since a late reaction (i.e. after faulty or quality violating services have been executed) might result in an inability to find a suitable recovery from that point, or a re-selected solution of lower quality compared to one that could be obtained when reacting to changes earlier; **(G4)** producing an optimal solution for any instantiated service re-selection process to avoid low-quality solutions; and **(G5)** keeping triggered adaptations transparent to the end user (i.e. eliminating execution interruption), since an interruption to the composite service execution could be highly undesirable, especially in the case of time-sensitive applications.

Existing approaches to adaptation usually achieve some of the above goals at the expense of others (see Fig. 1 for a comparison). Specifically, current approaches (e.g. [1–6]) can mostly be classified as *reactive*, performing corrective actions only after an exception has already occurred, thus lacking any ability to avoid erroneous behaviour or to improve performance, when possible. In addition, an interruption to the execution process is incurred until the corrective actions (usually through costly re-planning) are completed. Attempts to reduce such an interruption include applying fast heuristics (e.g. [4,5]) or pre-computing backup plans beforehand (e.g. [2,6]). While the former affects the solution quality, the constant changes of the service landscape may invalidate the latter, making the backups no longer applicable or poor-quality choices. Despite some recent efforts on *proactive* adaptation (e.g. [7–9]), they mostly focus on the early detection of exceptions, ignoring the actual adaptation process.

In response, this paper proposes an *early, efficient, and optimality-retaining* execution-time adaptive behaviour, capable of achieving *all* of the above goals during composite service execution. In particular, the paper makes the following contributions.

	G1	G2	G3	G4	G5
Current proactive attempts [9]	✓	✗	✓	✗	partially
Reaction on violation:					
execution-time heuristic re-selection [5]	✓	✗	✗	✗	partially
execution-time optimal re-selection [3]	✓	✗	✗	✓	✗
pre-computed backups [6]	partially	✗	✗	✗	✓

Fig. 1. Current adaptive composition approaches.

- *Parallel-to-execution reaction.* The adaptation process is instantiated at the *earliest* possibility during a component execution, so that the chance of completing the adaptation before the component execution terminates is maximised, and thus execution disruptions are minimised (**G5**). This early adaptation is instantiated for both corrective (**G1**) and optimising (**G2**) purposes. In particular, as opposed to existing approaches where the adaptation is mainly corrective, our adaptation is also triggered whenever an optimisation opportunity is identified (e.g. availability of new, better services), so that such an opportunity is exploited to improve the current solution. Moreover, since reaction to service changes is performed as soon as these occur in the environment, problems encountered in services scheduled for future execution are dealt with as early as possible (**G3**), before reaching erroneous execution points where recovery opportunities are of lower quality or not possible.
- *Change prioritisation.* Each service change encountered during execution goes through an assessment process to derive its priority regarding the situation at hand. Specifically, changes potentially affecting action points in the near future are handled urgently, while the adaptation to those of less importance is allowed to be carried out during the next component execution, without causing interruption (**G5**). Through such a novel and comprehensive analysis of changes, and the corresponding behaviour of the executing system, inefficiencies are avoided, unless necessary.
- *Light and optimality-retaining re-selection.* Our service re-selection algorithm *repairs* the affected part of an already existing search graph, without expensive recalculations from scratch, thus facilitating a fast adaptation (with almost no interruption) (**G5**). Moreover, the combination of services produced by this re-selection process is always the best possible, given the tasks already executed and the current environment state (**G4**).

This re-selection algorithm is built on our previous reactive service selection model [10], where efficient repair rules are introduced to incorporate reaction capabilities into static service selection [11] (which generates the original search graph at selection time). Here, however, we adjust this model so that repair remains possible at execution time, by assuming a *reverse* version of the search graph. This allows keeping the search graph at execution time up-to-date with the most recent environment state, via continuous light repair actions, thus facilitating fast reaction to any change during execution, especially in the critical case where the change concerns the task being invoked.

The paper is organised as follows. Section 2 discusses related work. The basic service selection model is summarised in Section 3, followed by a motivating example in Section 4. A classification of changes is introduced in Section 5, based on which the adaptive behaviour of the system is analysed in Section 6, and an efficient re-selection algorithm is outlined in Section 7. Sections 8 and 9 provide a theoretical and empirical analysis, respectively, and Section 10 concludes the paper.

2. Related work

Quality-based service selection has gained much attention from others. Like us, Yu et al. [12] and Li et al. [13] model it as a multi-constrained optimal path problem, and present heuristic algorithms to improve efficiency. In contrast, Canfora et al. [14] take a genetic algorithm approach. However, neither addresses adaptation to changes in a dynamic world.

To address the volatility of service environment, some efforts are aimed at fault avoidance, introducing preventive measures to reduce failures and quality deviations during execution, e.g. through redundant execution of services [15] or by providing accurate quality estimations [16]. Yet, since complete avoidance of execution-time exceptions is not possible, the ability to adapt to changes remains a critical requirement.

Many other efforts thus focus on achieving fault-tolerant behaviour to ensure that the system continues its intended execution, or at least terminates in a consistent state, in spite of the occurrence of failure or violation. In this regard, a number of approaches are concerned with incorporating exception handling mechanisms into the composition modelling language itself [17,18], allowing the designer (or user) to control recovery actions at execution time. Although effective for specific exception types (e.g. invalid input/output parameters), language-integrated adaptation may not be suitable for some other types (e.g. additions, deletions, or changes in quality values of services). This is because such environment changes are difficult to predict by the designer, and would result in an explosion of the exception handler complexities. Therefore, in this paper, adaptation is achieved at the middleware level.

Satisfying particular transactional patterns by the composite service has also been proposed in order to increase composition reliability and fault tolerance at execution time [19,20]. These efforts aim to minimise the risk for consumers by ensuring that the execution terminates in a consistent state even when failures occur, achieved through compensation policies allowing the effects of executed services to be undone. Such approaches, however, offer rather extreme and costly exception-handling capabilities, which may not be necessary in many situations, and are constrained to cooperative environments. Nevertheless, accounting for transactional properties can be considered an interesting extension to our approach.

A popular way of recovering from unexpected situations during execution (and the closest to our work) is by triggering re-planning actions in response. Some such efforts apply, during the re-planning stage, the same selection algorithm used to produce the initial solution, but incorporating the current execution status. For example, Zeng et al. [21] recalculate assignments for the non-executed part of a workflow each time a change occurs during execution by adopting Integer Programming. A re-planning triggering algorithm is introduced by Canfora et al. [1] to recalculate quality values of a composite service according to the new information at execution time (e.g., actual service qualities, or actual number of loop iterations), and if the new qualities differ considerably from previously estimated ones, execution is stopped and genetic-algorithm-based re-planning is triggered for remaining workflow tasks. A similar execution-time re-planning approach, but based on Integer Programming, is presented by Ardagna et al. [3]. Others introduce heuristic methods for the re-selection process to reduce its computational complexity. For example, Berbner et al. [4] use the H1_RELAX_IP heuristic, backtracking on the results of a relaxed integer program, to re-plan the remaining part of the workflow in a timely manner. Likewise, Lin et al. [5] propose a region-based heuristic re-selection algorithm, which iteratively expands the sub-process to be reconfigured until a satisfactory replacement is found. All these approaches can be categorised as *reactive*,

performing corrective actions only after faulty or quality-violating services are executed, thus ignoring emerging better opportunities, lacking the ability to prevent erroneous behaviour (even when such behaviour can be detected at an early stage), and causing an interruption to execution until re-selection is performed. That is, as opposed to our work, these approaches fail to achieve goals G2, G3, and G5.

In order to eliminate the undesired re-selection delay at execution time (goal G5), some approaches (e.g. [2,22,6]) suggest supporting the composite application with pre-computed backup services to ensure its continuous execution without any extra delay in the face of component failures. However, the problem with such approaches is that, due to the dynamic nature of services, the backups produced during selection may no longer remain optimal, satisfactory, or even available during execution. As a result, the execution could be faced with either a low-quality alternative, or a costly re-planning process to achieve a successful (or better) recovery.

Finally, although there are recent attempts towards achieving *proactive* adaptation (i.e. to prevent future failure or improve performance), these are still very limited and mainly focus on the change *detection* part, giving little or no consideration to the actual adaptation process. Proposed proactive change detection methods include applying performance prediction techniques [8,23,9], testing the behaviour of services using generated test cases [7,24], and subscribing to change requests with the registry [25]. Such detection efforts can be considered complementary to our work (which focuses instead on the latter change handling step). Like us, a few approaches also consider subsequent proactive adaptation actions (e.g. [8,9,26]), but these actions are mostly instantiated for corrective purposes, to prevent an anticipated problem, ignoring optimisation opportunities (G2). Furthermore, no proper modelling and management of the adaptation process, to avoid its interference with the application's execution (G5), is provided. These issues are addressed in our approach, achieving all goals (G1..G5), as summarised in Section 1 and detailed below.

3. Basic model

This section summarises the main components involved in the quality-based service selection problem, including our selection algorithm to solve this problem, originally introduced in [11]. See Fig. 2 for the notation used.


3.1. Planning knowledge model

The planning knowledge for a particular objective can be represented as a task hierarchy (T, tr, tf, tg) , where: T is a finite set of the tasks involved; tr is the root of the hierarchy (the goal task); tf is a functionality description function, assigning to each task $t \in T$ a semantic specification of its functional requirements; and finally tg is a task decomposition function, which maps each non-leaf task $t \in T$ to a set of directed acyclic graphs, each specifying a different way of decomposing t into finer-grained sub-tasks and their partial ordering constraints (execution order). An example *planning knowledge* for goal task *plan holiday* is shown in Fig. 3. In this Figure, task A can be either decomposed into sub-tasks B, C, and D (executed in sequence), or into sub-tasks G, C, and D (executed in sequence). Sub-task D is further decomposed into finer sub-tasks E and F (executed in sequence).

Note that task definition is kept generic to be applicable to a wide range of domains. It may refer, for example, to an operation signature (in terms of input and output parameters), to a resource specification, or simply to a term of an ontology agreed within a community. Moreover, different mechanisms are

Set	Description	Function	Description
<i>AD</i>	Services to be added to <i>rcnd</i>	<i>ion</i>	Node's index in a path
<i>AN</i>	Quality attribute names	<i>nodes</i>	Path's nodes
<i>AR</i>	Constrained attributes	<i>oi</i>	Task's optimal instances
<i>RM</i>	Services to be removed from <i>rcnd</i>	<i>rc</i>	Request's quality constraints
<i>S</i>	Available services	<i>rcnd</i>	Task's non-dominated services
<i>SPLN</i>	Request-based selection plans	<i>rt</i>	Requested task
<i>T</i>	Hierarchy tasks	<i>rw</i>	Request's quality weights
Function	Description	<i>sai</i>	Service appearing at an index
<i>abspln</i>	Requested task's abstract plans	<i>srdd</i>	Services dominated by a service
<i>cmp</i>	Requested task's composite services	<i>srds</i>	Services dominating a service
<i>cnd</i>	Task's candidate services	<i>su</i>	Service utility
<i>cu</i>	Comp. service utility	<i>sv</i>	Service's quality values
<i>cv</i>	Comp. service quality values	<i>tmn</i>	Task's min quality values
<i>en</i>	Path's last node	<i>tmx</i>	Task's max quality values
<i>es</i>	Instance's last service	<i>vldprd</i>	Task's valid predecessors
<i>ins</i>	Path's possible instances		

Fig. 2. Sets and functions used throughout the paper.

Fig. 3. Planning knowledge for *plan holiday* task.

possible for discovering suitable (candidate) services for each task: by consulting a central service repository storing service metadata (e.g., a semantic search over SAWSDL¹ descriptions of web services advertised in a UDDI registry); or by calling for service proposals over the network (e.g., using the contract net protocol [27]). We make no assumptions in our model about any specific technology or service discovery and matching mechanism, and leave this to the application domain, focusing instead on the generic problem of how to efficiently select and maintain the best combination of the available (discovered) services under the dynamism and uncertainty inherent in many such domains.

3.2. Service model

The space of available services can be defined as a tuple, (S, sf, sv) , where: S is the set of all available services; sf is a functionality description function, which assigns to each service $s \in S$ a semantic specification of its functionality, e.g. in OWL-S or WSDL-S; and finally sv is a quality of service (non-functional properties) specification function, which assigns to each service $s \in S$ its value for a quality attribute $a \in AN$ (AN is the set of all quality attributes).

Based on this, the candidate services for task $t \in T$, denoted $cnd(t) \subset S$, are those services $s \in S$ whose functional description, $sf(s)$, semantically matches the functional requirements of task t , $tf(t)$.

3.3. Request model

A composition request can be defined as a triple, (rt, rc, rw) . Task $rt \in T$ is the goal task to be accomplished. Function rc represents the QoS constraints imposed for task rt , and maps attribute $a \in AN$ to an upper or lower user-defined bound for its value, depending on the attribute *direction*. That is, $rc(a)$ is the minimum allowed value for attribute a if this attribute has an *increasing* direction (a higher value is better), or the maximum allowed value if attribute a 's direction is *decreasing* (a lower value is better). For simplicity, henceforth we assume that all quality attributes are decreasing. Note that $rc(a) = \text{undf}$ in case of no restrictions on the value of attribute a by the user. Finally, function rw specifies the user's preferences regarding different quality attributes, and assigns to each attribute $a \in AN$, a user-defined weighting factor $rw(a) \in [0, 1]$ reflecting its relative importance, s.t. $\sum_{a \in AN} rw(a) = 1$.

3.3.1. Request-based selection plans

Based on the planning knowledge model, multiple alternative *abstract plans* may be available for achieving the requested task rt . These plans, denoted $abspln(rt)$, correspond to all the possible expansions of the requested task, derived by recursively replacing task nodes with their decomposition graphs. For example, according to the planning knowledge hierarchy of Fig. 3, task A has five possible abstract plans: *plan1*: A; *plan2*: B–C–D; *plan3*: G–C–D; *plan4*: B–C–E–F; and *plan5*: G–C–E–F. Yet, not all these plans are necessarily interesting with respect to the user request at hand. That is, a plan whose available instances are all guaranteed to violate the quality constraints can be filtered out from the planning search space of the current request without affecting the ability to find an optimal solution. Formally, given a user request, the abstract plans to be considered for the selection process, denoted $SPLN$, are given as $SPLN = \{p \in abspln(rt) \mid \forall a \in AR, \text{aggr}_{t \in nodes(p)}(tmn(t, a)) \leq rc(a)\}$. Here, AR is the set of constrained quality attributes; $nodes(p)$ returns the task nodes of plan p ; $tmn(t, a)$ associates task t with the best (minimum value) offered for attribute a by this task's candidate services,

¹ <http://www.w3.org/TR/sawSDL/>.

i.e. $\min_{s \in \text{cnd}(t)} (sv(s, a))$; and **aggr** is some aggregation function that depends on the attribute considered. For example, possible aggregation functions for the quality attributes *execution time*, *reliability*, and *throughput* are the summation, multiplication, and minimum functions, respectively.

3.3.2. Request-based non dominated services

The set of alternative *composite services* for achieving the requested task, denoted $\text{cmp}(rt)$, is derived by *instantiating* plans SPLN (i.e. replacing the task nodes in each plan $p \in \text{SPLN}$ with a particular combination of their candidate services). With the increasing number of services per task, the number of such alternative compositions $|\text{cmp}(rt)|$ can be exponential. However, this number could be reduced considerably by filtering out from the candidate space of each task, all the services uninteresting for the current request. Such uninteresting services are those *request-based dominated* by another candidate service for the same task, with a service $s_j \in \text{cnd}(t)$ request-based dominating (r-dm) service $s_i \in \text{cnd}(t)$ **iff** s_i is worse than s_j regarding all the *constrained* quality attributes AR and the overall utility value su , i.e.

$$[\forall a \in AR, sv(s_i, a) \geq sv(s_j, a)] \wedge [su(t, s_i) \leq su(t, s_j)] \\ \wedge [\exists a \in AR, (sv(s_i, a) > sv(s_j, a)) \vee (su(t, s_i) < su(t, s_j))].$$

Here function $su(t, s) \in [0, 1]$ returns the overall utility of service $s \in \text{cnd}(t)$ regarding the user's request, s.t. $su(t, s) = \sum_{a \in AN} (rw(a) * \frac{tmx(t, a) - sv(s, a)}{tmx(t, a) - tmn(t, a)})$, where $tmx(t, a)$ is the maximum (and $tmn(t, a)$ the minimum) value offered for attribute a by task t 's candidate services.

Request-based dominated services are not potential candidates for the optimal solution, and thus can be ignored when instantiating plans SPLN .

3.4. Service selection problem

The service selection problem involves finding the best composite service to achieve the requested task, that both satisfies the user's imposed quality constraints and maximises the overall utility with respect to user-defined quality weights.

The value offered by a composite service $cs \in \text{cmp}(rt)$ for a particular quality attribute a , $cv(cs, a)$, is some aggregation **aggr** of the corresponding quality values for the component services, where **aggr** depends on the attribute considered. Based on this, the set of *satisfactory* composite services for the user's request, can be defined as $\text{SCS} = \{cs \in \text{cmp}(rt) \mid \forall a \in AN, (rc(a) \neq \text{undf}) \Rightarrow (cv(cs, a) \leq rc(a))\}$.

The solution composite service cs_{sol} for the user request is that satisfying: $cs_{sol} \in \text{SCS}$ such that $cu(cs_{sol}) = \max_{cs \in \text{SCS}} (cu(cs))$, where function $cu(cs) \in [0, 1]$ represents the overall utility of composite service cs , s.t. $cu(cs) = \sum_a (rw(a) * \frac{rmx(a) - cv(cs, a)}{rmx(a) - rmn(a)})$. Here, $rmx(a)$ returns the maximum (and $rmn(a)$ the minimum) value offered for attribute a by the requested task's actual plans (these maximum/minimum values can be estimated by aggregating for each abstract plan the tmx/tmn values of its tasks, and then calculating the maximum/minimum of these aggregated values).

3.5. Service selection algorithm

We model the service selection problem as a multi-constrained optimal path selection problem in a directed graph, called the *plan paths graph* (V_{PK}, E_{PK}) , where each path corresponds to an alternative abstract plan for achieving the requested task. We assume all abstract plans have a sequential structure (other structures can be transformed to the sequential structure using existing techniques [28]). Fig. 4 provides the plan paths graph for




Fig. 4. Plan paths graph for plan holiday task.

Node	A	B	C	D	E	F	G
<i>rcnd</i>	\emptyset	$s_{B1}(20, 30)$	$s_{C1}(15, 50)$	\emptyset	$s_{E1}(27, 5)$	$s_{F1}(30, 10)$	$s_{G1}(86, 5)$
(ex, pr)		$s_{B2}(30, 12)$	$s_{C2}(30, 30)$		$s_{E2}(15, 20)$	$s_{F2}(20, 40)$	
<i>vldprd</i>	\emptyset	S	SB	\emptyset	SBC	$SBCE$	\emptyset

Fig. 5. Request-based non-dominated services and valid predecessors for the nodes of Fig. 4.

the planning knowledge of Fig. 3. Note that this selection model views each abstract plan as an individual plan (with potential node overlap with other plans), and thus remains valid even if the planning knowledge is not hierarchical.

Based on the multi-constrained Bellman–Ford algorithm [29], our service selection algorithm is as follows. Each node v in the plan paths graph stores the optimal instances, denoted as $oi(v, p_v)$, for each path $p_v + v$ discovered so far from the start node to v (an instance of a path is a possible replacement of its task nodes with candidate services). In order to maximise utility, the concept of optimal paths in the original Bellman–Ford algorithm is updated so that an instance of path p is considered optimal if no other possible instance of the same path has both better values for all the constrained attributes and better utility. Moreover, to reduce the number of optimal instances, only those satisfying the quality constraints are maintained in each node. After traversing all graph nodes in topological order, the solution is the optimal composite service that has the best utility at the destination node.

In order to ensure that only plans SPLN are considered during selection, each node v in the plan paths graph is associated with the set of its valid predecessors $vldprd(v)$, which can be defined as follows: given a path $p_v + v$ from the start node to v , path p_v is considered a valid predecessor of node v if there exists at least one path p_i from v to the destination node, such that $p_v + p_i$ is a satisfactory abstract plan, i.e. $p_v + p_i \in \text{SPLN}$. Based on this, when processing an edge (u, v) , only the optimal composite services stored in u that are instances of v 's valid predecessors are considered. More details on this selection algorithm can be found in [11].

4. Example of service changes

Consider an example in which the user has issued a request to achieve task A, and is interested in minimising price (pr) while satisfying the constraint that execution time (ex) is less than 100. The plan paths graph for the requested task, and the request-based non-dominated candidate services of the sub-tasks involved are depicted in Figs. 4 and 5, respectively. Based on this, the set of plans to be considered for selection $\text{SPLN} = \{BCEF\}$ (plans A, BCD, GCD, and GCEF are excluded since tasks A and D do not have any available services, and all the available instances of plan GCEF violate the execution time constraint). Given set SPLN , the valid predecessors of the nodes are as shown in Fig. 5. The optimal solution for the user is instance $s_{B1}s_{C2}s_{E2}s_{F1}$ (ex:95, pr:90) which has the lowest price. In what follows, we give three example scenarios of service changes during the execution of the selected composite service $s_{B1}s_{C2}s_{E2}s_{F1}$.

Scenario 1. While executing service s_{B1} of composite service $s_{B1}S_{C2}S_{E2}S_{F1}$, a new service s_{D1} (ex:40, pr:10) joins the candidate services of task D . As a result, plan BCD is added to set $SPLN$, and two additional instances of this plan are satisfactory from this point, of which composite service $s_{B1}S_{C2}S_{D1}$ (ex:90, pr:70) is better than the selected composition $s_{B1}S_{C2}S_{E2}S_{F1}$ for both price and execution time.

Scenario 2. While executing service s_{B1} of composite service $s_{B1}S_{C2}S_{E2}S_{F1}$, service s_{C2} becomes unavailable. Here, simply replacing s_{C2} with s_{C1} will result in composition $s_{B1}S_{C1}S_{E2}S_{F1}$ (ex:80, pr:110), which is not optimal regarding price. Hence, both s_{C2} and s_{E2} should be substituted in this case in order to obtain the new optimal satisfactory solution, which is service $s_{B1}S_{C1}S_{E1}S_{F1}$ (ex:92, pr:95).

Scenario 3. While executing service s_{B1} of composite service $s_{B1}S_{C2}S_{E2}S_{F1}$, service s_{C2} changes its quality values to (ex:50, pr:20). From this point, the selected composite service is no longer satisfactory, and the new optimal satisfactory one is $s_{B1}S_{C1}S_{E1}S_{F1}$ (ex:92, pr:95).

5. Service change categorisation

As illustrated above, a change to the service landscape during execution may cause corresponding changes in the optimal composite services possible from that point (and potentially affecting the best solution), thus necessitating their recalculation in response (which we refer to as the re-selection process). Generally, the importance and urgency of responding to an encountered service change, i.e. triggering the re-selection process, vary depending on whether this change affects the non-dominated services of the respective task and other factors. Based on this, we propose categorising execution-time service changes into changes not to be considered and changes to be considered. These categories are detailed next after modelling the effect of a service change on the request-based non-dominated services of the task affected. In what follows, α_o and α_n represent α before and after a change occurrence.

5.1. The effect on non-dominated services

A change to the available services of task $t_{ch} \in V_{PK}$, might affect this task's set of request-based non-dominated services $rcnd(t_{ch})$, causing the addition of new services AD while removing existing ones RM , i.e. $rcnd_n(t_{ch}) = (rcnd_o(t_{ch}) \setminus RM) \cup AD$. The definition of sets AD and RM depends on the type of change that occurred.

In particular, where a new service s_n joins the candidate services of task t_{ch} , services AD to be added to $rcnd_o(t_{ch})$ is service s_n if it is not request-based dominated by an existing service in $rcnd_o(t_{ch})$, while services RM to be removed from $rcnd_o(t_{ch})$ are those existing services in $rcnd_o(t_{ch})$ that are request-based dominated by s_n .

Where a candidate service s_o of task t_{ch} becomes unavailable, services RM to be removed from $rcnd_o(t_{ch})$ is service s_o if it is a member of $rcnd_o(t_{ch})$, while services AD to be added to $rcnd_o(t_{ch})$ are task t_{ch} 's candidate services that are request-based dominated by s_o , and which, as a result of eliminating s_o , become non-dominated according to current request.

Where a candidate service s_o of task t_{ch} changes its quality values, with s_{ch} denoting this service after the change, the case can be treated as a deletion of s_o , followed by a subsequent addition of s_{ch} .

Formal definitions of sets AD and RM per each change type can be found in [30].

5.2. Changes not to be considered

A change to the available services of task t_{ch} while executing task t_{inv} (the task currently invoked of the selected solution) need not be considered, i.e. does not trigger the re-selection process, **iff** one of the following is satisfied.

- Task t_{ch} is already executed. That is, $t_{ch} = t_{inv}$ or t_{ch} appears *before* t_{inv} according to the topological order of the plan paths graph. Note that we concentrate, in our work, on a forward recovery mechanism, where only the tasks that are not yet executed can be re-assigned to other services or replaced with other tasks from an alternative plan, while those already executed are considered final with their assignments being unchangeable. This is a plausible assumption for some types of services (examples of which are the learning objects evaluated in Section 9, where it does not seem reasonable to make the user repeat the learning process for an already acquired concept). Extending our mechanism to allow for the possibility of rolling back the execution to a previous point in time is the focus of future work (see Section 10).
- The request-based non-dominated services of task t_{ch} are not affected by the change, i.e. $AD = RM = \emptyset$. Since non-dominated services are the only candidates for the optimal solution, there is no need to respond to this change.
- Task t_{ch} is not part of the plan being executed (p_{sel}) and does not belong to any satisfactory plan after the change. That is, $(t_{ch} \notin nodes(p_{sel})) \wedge (vldprd_n(t_{ch}) = \emptyset)$. In this case, there is no need to respond to the change even if it impacts the request-based non-dominated services of task t_{ch} , since this task is not part of any plan that will lead to a satisfactory solution for the current request.

5.3. Changes to be considered

A change to the available services of task t_{ch} while executing task t_{inv} needs to be considered, i.e. triggers the re-selection process, **iff all** of the following are satisfied: task t_{ch} is not executed yet, i.e. t_{ch} appears *after* t_{inv} according to the topological order of the plan paths graph; t_{ch} belongs to the plan being executed (p_{sel}) or belongs to at least one satisfactory plan after the change, i.e. $(t_{ch} \in nodes(p_{sel})) \vee (vldprd_n(t_{ch}) \neq \emptyset)$; and the request-based non-dominated services of task t_{ch} are affected by the change, i.e. $(AD \neq \emptyset) \vee (RM \neq \emptyset)$. Changes to be considered are further divided into non-affecting changes and affecting changes, as detailed next.

5.3.1. Non-affecting changes

A change is non-affecting if it has an impact on the optimal composite services possible from that point, but does not affect the best solution (the need to respond to this category of change is justified in Section 7). Having no effect on the best solution, this category of change does not cause any delay to the execution process. In other words, the solution composite service can continue its execution even if the re-selection process in response to the change is still running. Generally, a change to be considered is regarded as non-affecting in the following cases: the deletion of a non-selected service (a service that is not part of the current best solution); and changes in the quality values of a non-selected service s_o (s_{ch} denotes this service after the change) such that $(s_o \text{ r-dm } s_{ch})$ **or** $((s_o \text{ is incomparable with } s_{ch}) \text{ and } (s_{ch} \notin AD))$.

5.3.2. Affecting changes

A change is affecting if it has an impact on the optimal composite services possible from that point, and could cause a modification to the best solution. This category is further divided into non-interrupting changes and interrupting changes.

Non-interrupting changes are those affecting changes, the reaction to which does not cause any interruption between service executions, since the next service to be executed can be identified without the need for re-selection to be completed. Specifically, an affecting change to the services of task t_{ch} is non-interrupting **iff** task t_{ch} is the next task to be executed according to the current solution, with service s_{sel} being the currently selected service for this task, and *one* of the following is satisfied: the change that occurred is the addition of a new service s_n such that s_n r-dm s_{sel} ; or the change is a modification in the quality values of service s_o (s_{ch} denotes this service after the change) such that s_{ch} r-dm s_{sel} (note that s_{sel} might be the service affected by the change, i.e. $s_o = s_{sel}$). Intuitively, responding to such changes will result in replacing service s_{sel} with service s_n (in the addition case), and with service s_{ch} (in the modification case). Hence, the next service can be anticipated without requiring interruption.

Interrupting changes are those affecting changes, the reaction to which might result in an interruption to the composite service execution. This is because the next service to be executed cannot be identified prior to performing re-selection, thus causing the execution process to stop until re-selection is completed. Specifically, an affecting change to the services of task t_{ch} is interrupting in the following cases. *Case 1*: the addition of a new service s_n such that *one* of following is satisfied: t_{ch} is not part of the plan being executed; t_{ch} belongs to the plan being executed and s_n is incomparable with s_{sel} (the currently selected service for task t_{ch}); or t_{ch} belongs to the plan being executed, but is not the next task in the execution sequence, and s_n r-dm s_{sel} . *Case 2*: the deletion of a selected service. *Case 3*: changes in the quality values of a non-selected service s_o (s_{ch} denotes this service after the change) such that *all* of the following are satisfied: $[s_{ch}$ r-dm $s_o] \vee [(s_{ch}$ is incomparable with $s_o) \wedge (s_{ch} \in AD)]$; and $[t_{ch}$ is not the next task to be executed] $\vee \neg[s_{ch}$ r-dm $s_{sel}]$, where s_{sel} is the currently selected service for task t_{ch} . *Case 4*: changes in the quality values of a selected service s_{sel} (s_{ch} denotes this service after the change) such that t_{ch} is not the next task to be executed or $\neg(s_{ch}$ r-dm $s_{sel})$. Note that all the three change scenarios in Section 4 are considered interrupting, satisfying Case 1, Case 2, and Case 4, respectively.

6. Adaptive execution behaviour

Delaying the re-selection process until a violating behaviour is invoked results in undesired effects at execution time. For instance, observing the unavailability of service s_{C2} in [Scenario 2](#) only when trying to invoke this service causes execution to stop until re-selection is performed. Similarly, in [Scenario 3](#), detecting the changes in the quality values of service s_{C2} only after its execution results in an unrecoverable situation, since no satisfactory solution can be found from this point (all the service combinations including services s_{B1} and s_{C2} violate the user's execution time constraint).

To tackle this, we propose an early, parallel-to-execution adaptive system behaviour, where adaptation to changes is performed as soon as these changes occur in the environment, concurrently with the execution of the current service, thus reducing the delay between service executions, and increasing the chance of a successful recovery. For instance, in [Scenario 2](#), re-selecting services for tasks C, E and F in response to the deletion of service s_{C2} can be achieved while executing service s_{B1} , without causing extra delay.

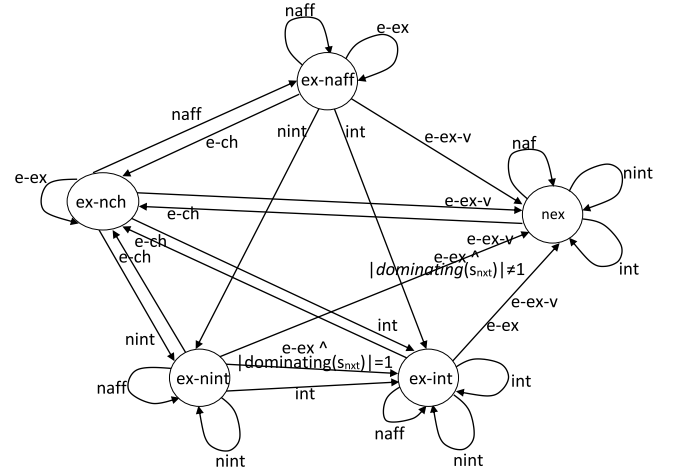


Fig. 6. Adaptive behaviour during execution.

Based on the change categories introduced, such an adaptive execution behaviour can be modelled using the finite state automaton in [Fig. 6](#), which consists of five states. States $ex - \alpha$ indicate that a component service of the best solution is currently running and, at the same time, the following is satisfied according to the value of α : when $\alpha = nch$, no re-selection is being performed by the system; when $\alpha = naff$, a re-selection is being performed in response to a set of non-affecting changes; when $\alpha = nint$, a re-selection is being performed in response to a set of changes including at least one non-interrupting change and no interrupting changes; and finally, when $\alpha = int$, a re-selection is being performed in response to a set of changes including at least one interrupting change. State nex indicates that the best solution execution is currently interrupted until re-selection is completed.

The behaviour of the system is interpreted as follows. The execution begins in state $ex-nch$, by invoking the first component service. With the occurrence of a change to be considered during a component execution, the system transitions to one of the states $ex-naff$, $ex-nint$, or $ex-int$, based on the change category, which could be a non-affecting change (event $naff$), an affecting and non-interrupting change (event $nint$), or an affecting and interrupting change (event int). The change category is identified with respect to the currently selected solution (we assume that the time required for this identification is negligible, especially when compared to re-selection time). The selected solution may be updated each time a re-selection is completed (event $e-ch$), causing the system to return to state $ex-nch$.

After the successful execution of a component service (event $e-ex$), the state into which the system transitions is determined based on its current state, as follows. If the system is in state $ex-nch$, i.e. no re-selection is being performed, the next service in the currently selected solution is invoked, without changing the state of the system. If the system is in state $ex-naff$, i.e. the re-selection being performed will not affect the currently selected solution, the next service in this solution is invoked, without changing the state of the system. In other words, the re-selection is carried on while executing the next service. If the system is in state $ex-int$, i.e. the next service to be executed cannot be identified before completing the re-selection being performed, the system transitions to state nex , and remains in this state until re-selection is completed and the next service to be invoked is determined. Finally, if the system is in state $ex-nint$, two cases are distinguished according to set $srds(s_{nxt})$, the set of services dominating the currently selected service s_{nxt} for the next task in the execution order, among the request-based non-dominated services of this task. *Case 1*: $|srds(s_{nxt})| = 1$, in which the next service to be executed can be estimated without the need for re-selection to

be completed. This service, $s_{nxt-new} \in srds(s_{nxt})$, is thus invoked without delaying execution, causing the system to transition to state *ex-int*. In other words, the re-selection is continued while executing service $s_{nxt-new}$, but is considered interrupting since the next service to invoke after service $s_{nxt-new}$ cannot be known prior to completing re-selection. Case 2: $|srds(s_{nxt})| \neq 1$, in which the next service to be executed cannot be determined before the re-selection is completed. Therefore, the execution process is interrupted by moving to state *nex* to continue the re-selection.

The case where the current component service delivers unexpected quality values (event *e-ex-v*) is considered an interrupting change, and thus also causes the system to enter state *nex*, regardless of its current state.

Example. Consider **Scenario 2** of Section 4, with the initial solution $s_{B1}S_{C2}S_{E2}S_{F1}$. Invoking s_{B1} initiates the adaptive execution behaviour at state *ex-nch*. Since the deletion of s_{C2} while executing s_{B1} is an interrupting change, it triggers the transition of the system to state *ex-int* to indicate a running re-selection. Once the re-selection is completed, the system goes back to state *ex-nch*, updating the solution to $s_{B1}S_{C1}S_{E1}S_{F1}$.

7. Efficient service reselection

The adaptive behaviour proposed above triggers re-selection in response to changes in parallel with execution, in order to avoid extra delays. However, a costly re-selection process could still cause an interruption to execution, especially if the change is only discovered at a late stage or (in the worst case) at the end of the current component execution. In response, we introduce here a light re-selection approach, applying efficient repair rules to an already existing search graph (the graph produced by the initial selection process), without expensive recalculations from scratch. The idea is to apply the selection algorithm (of Section 3.5) prior to execution, in order to generate the search graph (i.e. generate the optimal instances for each task node), and to select the initial solution. The search graph is then kept valid during execution by continuously adjusting it with respect to the environment state (which justifies the need to account for non-affecting changes). Maintaining the graph validity ensures that, whenever any change occurs (especially a critical, affecting one), only a minimal number of modifications to the affected part of the graph are required in response, thus increasing the chance that the adaptation to the change terminates before the end of the current component execution. Next, we first introduce the search graph enabling efficient execution-time adaptivity, followed by the graph repair rules in response to a change (i.e. the re-selection algorithm).

7.1. Execution-time search graph

In the simplest case, with no changes encountered, the validity of the search graph should be maintained against the execution progress of the selected solution. This, however, could be costly to achieve with the forward version of the plan paths graph, where the task nodes are processed (by the selection algorithm) according to their execution order. To illustrate, consider the example of Section 4, with the forward plan paths graph in Fig. 4, and the initial solution $s_{B1}S_{C2}S_{E2}S_{F1}$. Once service s_{B1} is invoked, the optimal instances recorded at the remaining, non-executed task nodes (i.e. tasks nodes C, D, E, and F) are no longer valid. This is because these instances (which correspond to paths beginning with node B) account for service s_{B2} as a possible candidate for executing task B, no longer the case after s_{B1} 's execution.

To tackle this, we apply the selection algorithm on the *reverse* version of the plan paths graph, generated by reversing the direction of edges in the original plan paths graph (i.e. the




Fig. 7. Reverse plan paths graph for *plan holiday* task.

start node of the reverse plan paths graph is the end node of the original one). For example, the reverse graph for the plan paths graph of Fig. 4 is provided in Fig. 7. Such modified selection produces the same best solution, while maintaining the validity of optimal instances at the non-executed task nodes as the execution advances, due to the reverse processing order of nodes. For instance, in our example, when selection is performed on the reverse graph of Fig. 7, executing service s_{B1} does not affect the optimal instances of task node C which, in this case, records instances of paths *DC* and *FEC*. The same holds for the other, non-executed task nodes D, E and F. Only the optimal instances of the node executed, B, are affected, and would need to be re-computed if a change occurs (see Section 7.2).

Based on this, each time a new component service, s_{inv} , of the selected solution, $ins_{ex} + s_{inv} + ins_{uex}$, is invoked (where ins_{ex} and ins_{uex} correspond to the executed and non-executed parts of this solution, respectively), only the following adjustments are required to the search graph: changing the destination node to task t_{inv} (the task being executed), with its request-based non-dominated services being set to instance $ins_{ex} + s_{inv}$; and updating the valid predecessors of nodes through adjusting the selection plans *SPLN*, such that all the plans not beginning with the already executed path are removed from set *SPLN*, since these are invalid from this point. In our example, reflecting the execution of service s_{B1} , on the reverse search graph, would only involve changing the destination node to node B, and adjusting $rcnd(B)$ to $rcnd(B) = \{s_{B1}\}$. Here, no modification is required to set *SPLN* (and consequently the valid predecessors) since only plan *BCEF*, currently under execution, is included in this set.

7.2. Search graph repair rules

To account for a service change to be considered at task node t_{ch} , while executing service s_{inv} (of task t_{inv}), we apply only the necessary updates to the optimal instances at the nodes of the reverse search graph, without re-computing these instances from scratch. This is achieved by associating each valid predecessor p_u , of each node u , with three mutable sets capturing the updates required: an *additional services* set $as(u, p_u) \subset rcnd(u)$ specifying what services of node u need to be joined with path p_u 's optimal instances when updating the optimal instances of path $p_u + u$; an *additional instances* set $ai(u, p_u, i \in \mathbb{Z}^+) \subset S$ specifying what optimal instances of path p_u need to be joined with node u 's services when updating the optimal instances of path $p_u + u$ (i.e. $s \in ai(u, p_u, i)$ indicates that, of the additional optimal instances ins of path p_u to be combined with node u 's services, are those including service s at position i); and a *domination check* set $dc(u, p_u) \subset ins(p_u + u)$ specifying what optimal instances of path $p_u + u$ become unavailable, thus, when updating the optimal instances of path $p_u + u$, all its instances previously dominated by at least one instance in $dc(u, p_u)$ should be checked for optimality.

Algorithm 1 summarises the repair process of the optimal instances of path $p_u + u$ at node u , according to above semantics. Procedure *check-instance-optimality* (ins, u, p_u) assesses the optimality of instance ins against those already recorded at node u for path $p_u + u$. Note that $as(u, p_u) = ai(u, p_u, i \in \mathbb{Z}^+) = dc(u, p_u) = \emptyset$, when no modifications to the optimal instances of path $p_u + u$ are

Algorithm 1 Repair-optimal-instances(u, p_u)

```

1:  $v \leftarrow en(p_u)$ 
2:  $p_v \leftarrow p_u - v$ 
3: if  $dc(u, p_u) \neq \emptyset$  then
4:   for each  $s \in rcnd(u) \setminus as(u, p_u)$  do
5:     for each  $ins_v \in oi(v, p_v)$  s.t.  $\forall i \in \mathbb{Z}^+, sai(ins_v, i) \notin ai(u, p_u, i)$  do
6:       if  $\exists ins \in dc(u, p_u), ins \text{ r-dm } ins_v + s$  then
7:         check-instance-optimality( $ins_v + s, u, p_u$ )
8:   if  $\exists i \in \mathbb{Z}^+, ai(u, p_u, i) \neq \emptyset$  then
9:     for each  $ins_v \in oi(v, p_v)$  s.t.  $\exists i \in \mathbb{Z}^+, sai(ins_v, i) \in ai(u, p_u, i)$  do
10:      for each  $s \in rcnd(u) \setminus as(u, p_u)$  do
11:        if  $\forall a \in AR, cv(ins_v + s, a)$  is-better-than  $rc(a)$  then
12:          check-instance-optimality( $ins_v + s, u, p_u$ )
13:   if  $as(u, p_u) \neq \emptyset$  then
14:     for each  $s \in as(u, p_u)$  do
15:       for each  $ins_v \in oi(v, p_v)$  do
16:         if  $\forall a \in AR, cv(ins_v + s, a)$  is-better-than  $rc(a)$  then
17:           check-instance-optimality( $ins_v + s, u, p_u$ )
18:  $as(u, p_u) \leftarrow \emptyset; ai(u, p_u, i \in \mathbb{Z}^+) \leftarrow \emptyset; dc(u, p_u) \leftarrow \emptyset$ 

```

Procedure 2 Check-instance-optimality(ins, u, p_u)

```

1:  $optml \leftarrow 1$ 
2: for each optimal instance  $ins_u \in oi(u, p_u)$  do
3:   if  $ins_u \text{ r-dm } ins$  then
4:      $optml \leftarrow 0$ 
5:   break
6: else if  $ins \text{ r-dm } ins_u$  then
7:   remove  $ins_u$  from the instances at  $oi(u, p_u)$ 
8: if  $optml=1$  then
9:   add  $ins$  to the instances  $oi(u, p_u)$ 

```

Node	F	E	D	C	B	A	G
<i>vldprd</i>	S	SF	\emptyset	SFE	SFEC	\emptyset	\emptyset
<i>as</i>	\emptyset	$\{s_{E3}\}$	—	\emptyset	$\{s_{B1}\}$	—	—
<i>ai</i>	\emptyset	\emptyset	—	$\{s_{E3}, \text{index} = 2\}$	\emptyset	—	—
<i>dc</i>	\emptyset	\emptyset	—	\emptyset	\emptyset	—	—

Fig. 8. Sets *as*, *ai*, and *dc* in the running example.

required. Once the repair for the nodes is completed, the new best solution is the one with the highest utility among the adjusted optimal instances at node t_{inv} (the current end node).

The instantiation of sets *as*, *ai*, and *dc*, for each valid predecessor, depends on the change type, and full details of such an instantiation can be found in [30].

Example. Suppose that while service s_{B1} of the initial solution $S_{B1}S_{C2}S_{E2}S_{F1}$ of Section 4 is executed, service $s_{E3}(10, 50)$ joins the services of node E. This change is a change to be considered with $AD = \{s_{E3}\}$, $RM = \{\emptyset\}$, and no effect on the valid predecessors of nodes. Given the search graph of Fig. 7, the instantiation of sets *as*, *ai*, and *dc*, for the valid predecessors, in response to this change, is provided in Fig. 8.

8. Analytical study

This section analyses the time complexity of the proposed repair-based re-selection algorithm, and compares it against re-selection from scratch (also assumed to be applied on the reverse version of the plan paths graph for ease of comparison). We focus here on the step of optimal instances modification in response to a change since, when compared to this step, the time required for the other change handling steps (e.g. updating the request-based non-dominated services of affected nodes and identifying the change category) is negligible.

Since the valid predecessors at each node are processed independently, the analysis assumes for simplicity one valid predecessor per node (the specific case of one abstract plan).

This still allows demonstration of the efficiency gain achieved by our approach for any affected valid predecessor, and can easily be generalised to handle the case of multiple valid predecessors (i.e. multiple abstract plans). Note that, in such a general case, our approach achieves further time reduction due to reprocessing only the *affected* valid predecessors per each node, compared to the re-selection from scratch which reprocesses *all* the valid predecessors.

Next, the time of the pre-execution selection algorithm is analysed in order to provide the basis for the subsequent analysis of re-selection approaches.

8.1. Selection algorithm

Consider a sequential abstract plan comprising k tasks, $v_1 v_2 \dots v_k$, each with n available candidate services. To select the best solution (the best instance of path $v_1 v_2 \dots v_k$), each node $v_{i>1}$ records the optimal instances of path $v_1 v_2 \dots v_i$, denoted $oi(v_i)$. Hence, selection time $\tau(sel)$ is:

$$\tau(sel) = \sum_{i=2}^k \tau(oi(v_i)). \quad (1)$$

The time required to calculate $oi(v_i)$, $\tau(oi(v_i))$, depends on the sizes of $oi(v_{i-1})$ and $rcnd(v_i)$, which can be defined in terms of the following pruning rates: $spr_i \in [0, 1]$, denoting the percentage of candidate services pruned per task node v_i prior to selection, i.e. $|rcnd(v_i)| = n \times sr_i$, where $sr_i = 1 - spr_i$; and $cpr_i \in [0, 1]$, denoting the percentage of instances pruned per path $v_1 v_2 \dots v_i$ when computing the optimal instances at node v_i , i.e. $|oi(v_i)| = |oi(v_{i-1})| \times |rcnd(v_i)| \times cr_i$, where $cr_i = 1 - cpr_i$. Since $|oi(v_1)| = |rcnd(v_1)| = n \times sr_1$,

$$|oi(v_i)| = n^i \times \prod_{m=1}^i sr_m \times \prod_{m=2}^i cr_m. \quad (2)$$

Based on this, $\tau(oi(v_{i>1}))$ is given as follows:

$$\begin{aligned} \tau(oi(v_i)) &\text{ is } O(|oi(v_{i-1})| \times |rcnd(v_i)|^2) \\ &\text{ is } O\left(\left(n^{i-1} \times \prod_{m=1}^{i-1} sr_m \times \prod_{m=2}^{i-1} cr_m \times n \times sr_i\right)^2\right) \\ &\text{ is } O\left(n^{2i} \times \prod_{m=1}^i sr_m^2 \times \prod_{m=2}^{i-1} cr_m^2\right). \end{aligned} \quad (3)$$

From Eqs. (1) and (3), the time complexity of selection $\tau(sel)$ is $O(n^{2k} \times \prod_{m=1}^k sr_m^2 \times \prod_{m=2}^{k-1} cr_m^2)$. Note here that $k \ll n$.

Assuming for simplicity that $\forall i, sr_i = cr_i = r$, $\tau(sel)$ is $O(n^{2k} \times r^{4k-4})$. Hence, our service selection achieves a time complexity of $O(n^\alpha)$ if rate $r = \sqrt[4k-4]{n^{\alpha-2k}}$. For example, in order for selection to be of linear time complexity, i.e. $\alpha = 1$, when $n = 100$ and $k = 5$, rate r should be: $r = \sqrt[16]{100^{-9}} = 0.08$. That is, the pruning rate $(1-r)$ should be at least 92%.

8.2. Reselection from scratch

The re-selection from scratch approach recalculates the optimal instances of all non-executed nodes from scratch, in response to a change to be considered at node v_{ch} while executing node v_{inv} . Thus, its time complexity, $\tau^s(resel)$, is:

$$\tau^s(resel) = \sum_{i=1}^{inv-1} \tau^s(oi_n(v_i)). \quad (4)$$

Here, $\tau^s(o_{i_n}(v_i))$ is the time required for *recomputing* the optimal instances at node v_i , given as (see Eq. (3)):

$$\tau^s(o_{i_n}(v_{i < ch})) \text{ is } O\left(n^{2i} \times \prod_{m=1}^i sr_m^2 \times \prod_{m=2}^{i-1} cr_m^2\right)$$

$$\tau^s(o_{i_n}(v_{i \geq ch})) \text{ is } O\left(n^{2(i-1)} \times n'^2 \times \prod_{m=1}^i sr_m^2 \times \prod_{m=2}^{i-1} cr_m^2\right)$$

where $n' = |cnd_n(v_{ch})|$, i.e. $n' = n + 1$ in case of service addition; $n' = n - 1$ in case of service deletion; and $n' = n$ in case of changes in service qualities. Based on this and Eq. (4), we can conclude that ($inv \leq k$):

$$\tau^s(resel) \text{ is } O\left(n^{2(inv-1)} \times \prod_{m=1}^{inv-1} sr_m^2 \times \prod_{m=2}^{inv-2} cr_m^2\right). \quad (5)$$

8.3. Repair-based reselection

The proposed repair-based re-selection approach only makes the updates necessary to the affected optimal instances, in response to a change to be considered at node v_{ch} while executing node v_{inv} , without recalculating those instances from scratch. Thus, its time complexity, $\tau^r(resel)$, is:

$$\tau^r(resel) = \sum_{i=ch}^{inv-1} \tau^r(o_{i_n}(v_i)). \quad (6)$$

Here, $\tau^r(o_{i_n}(v_i))$ is the time required for *modifying* the optimal instances at node v_i . The modification depends on the type of change that has occurred, and is analysed next for the addition case, i.e. addition of a request-based non-dominated service s_n at node v_{ch} (the deletion and quality changes cases can be analysed similarly).

For node v_{ch} , the modification involves combining the optimal instances of node v_{ch-1} with service s_n , and then checking the optimality of the resulting combinations against those already recorded at node v_{ch} , i.e. $\tau^r(o_{i_n}(v_{ch}))$:

$$\text{is } O(|o_{i_o}(v_{ch-1})| \times |o_{i_o}(v_{ch})|)$$

$$\text{is } O\left(n^{ch-1} \times \prod_{m=1}^{ch-1} sr_m \times \prod_{m=2}^{ch-1} cr_m \times n^{ch} \times \prod_{m=1}^{ch} sr_m \times \prod_{m=2}^{ch} cr_m\right)$$

$$\text{is } O\left(n^{2ch-1} \times \prod_{m=1}^{ch-1} sr_m^2 \times \prod_{m=2}^{ch-1} cr_m^2 \times sr_{ch} \times cr_{ch}\right). \quad (7)$$

Similarly, for node $v_{i > ch}$, updating $o_{i_n}(v_i)$ involves checking the optimality of the newly available instances (obtained by joining the optimal instances containing service s_n at node v_{i-1} , $o_{i_n}(v_{i-1})^{s_n}$, with node v_i 's services), against those already computed at node v_i , i.e. $\tau^r(o_{i_n}(v_{i > ch}))$:

$$\text{is } O(|o_{i_n}(v_{i-1})^{s_n}| \times |rcnd(v_i)| \times |o_{i_n}(v_i)|)$$

$$\text{is } O\left(\frac{n^{i-2} \times n' \times \prod_{m=1}^{i-1} sr_m \times \prod_{m=2}^{i-1} cr_m}{n' \times sr_{ch}} \times n \times sr_i \times n^{i-1}\right.$$

$$\left. \times n' \times \prod_{m=1}^i sr_m \times \prod_{m=2}^i cr_m\right)$$

$$\text{is } O\left(n^{2(i-1)} \times n' \times \prod_{\substack{m=1 \\ m \neq ch}}^i sr_m^2 \times \prod_{m=2}^{i-1} cr_m^2 \times sr_{ch} \times cr_i\right). \quad (8)$$

Here, $n' = n + 1$, and $|o_{i_n}(v_i)^{s_n}| = \frac{|o_{i_n}(v_i)|}{|rcnd(v_{ch})|}$. From Eqs. (6)–(8), we conclude that $\tau^r(resel)$ is:

$$O\left(n^{2inv-3} \times \prod_{\substack{m=1 \\ m \neq ch}}^{inv-1} sr_m^2 \times \prod_{m=2}^{inv-2} cr_m^2 \times sr_{ch} \times cr_{inv-1}\right). \quad (9)$$

This is applicable as long as $ch < inv$ (the node affected by the change is not the node being executed). Yet, when $ch = inv$ (the invoked service delivers unexpected qualities), reselection only involves recombining the optimal instances already recorded at node v_{inv-1} with the modified invoked instance, and thus $\tau^r(resel)$ is $O(|o_{i_o}(v_{inv-1})|)$, i.e.

$$\tau^r(resel) \text{ is } O\left(n^{inv-1} \times \prod_{m=1}^{inv-1} sr_m \times \prod_{m=2}^{inv-1} cr_m\right). \quad (10)$$

8.4. Comparison

To analyse the efficiency gain achieved by the proposed repair-based re-selection (compared to reselection from scratch), we make the simplifying assumption that $\forall i, sr_i = cr_i = r$. As a result, comparing time complexities τ^s and τ^r , leads to:

$$ch < inv : \frac{\tau^s(resel)}{\tau^r(resel)} = \frac{n^{2inv-2} \times r^{4inv-8}}{n^{2inv-3} \times r^{4inv-8}} = n$$

$$ch = inv : \frac{\tau^s(resel)}{\tau^r(resel)} = \frac{n^{2inv-2} \times r^{4inv-8}}{n^{inv-1} \times r^{2inv-3}} = n^{inv-1} \times r^{2inv-5}.$$

In other words, when the change occurs at a non-executed node, the proposed approach reduces reselection time by a factor of n . The reduction factor further increases to $n^{inv-1} \times r^{2inv-5}$ in the case where the change affects the node being executed (i.e. cannot be anticipated in advance).

9. Empirical study

The goal of this section is to assess the efficiency of our repair-based re-selection algorithm, the gain in utility by responding to changes ahead of time, and the reduction in execution interruption by performing re-selection without interfering with the execution process (unless necessary).

9.1. Experimental setup

We perform the evaluation in domain of learning object composition [31], where the goal is to fulfil a particular learning objective by automatically compositing existing reusable learning objects (the candidate services in our model) into a respective course, taking learner (user) preferences and constraints into consideration (full details on this case study can be found elsewhere [30]). Learning objects (LOs) are published through learning object repositories, where hundreds of learning objects with different properties can be available for each concept, and are usually heterogeneous in their granularities, i.e. they can range from a single image to a whole module. Thousands of new learning objects are made available every day, while existing learning objects can be updated or become unavailable at any




Fig. 9. Reselection time in response to a random change.

time. Such changes in the repository can occur during the delivery (execution) of a selected course (a composition of learning objects), thus possibly necessitating the re-selection of learning objects for concepts not yet presented, in order to guarantee the most suitable learning experience for the user.

The metadata elements (quality of service properties) of 36 956 learning objects were collected from 10 different repositories, using the OAI-PMH² protocol. Of those elements, we selected the following for the global-level constraints and utility function: interactivity type, semantic density, difficulty, typical learning time, size, and cost. The planning knowledge adopted is a hierarchical representation of the *Algorithms and Data Structure* domain, where the tasks correspond to domain concepts. It comprises 3 hierarchical levels, with up to 11 concepts (tasks) per abstract plan in the search graph.

Learning objects are examples of services with long execution durations $execTime$ (corresponding to learning time). For such long-running services, where $execTime \gg rslScratch$ ($rslScratch$ is the time anticipated to perform the most costly re-selection from scratch, for the non-executed tasks), the proposed continuous adaptive behaviour may cause an unnecessary overhead from the composite service provider perspective. This is because, when a service execution spans a long period of time $[t_s, t_e]$, it becomes unnecessary to continuously react to changes for this entire duration. Instead, all changes could be ignored until $t_{crit} = t_e - rslScratch$, at which point re-selection from scratch for the remaining tasks should be instantiated to restore a valid instance of the search graph with respect to the new environment state. From this point, the execution should continue in the proposed light repair-based manner, efficiently accommodating all the changes occurring during the critical interval $[t_{crit}, t_e]$ (the interval signalling the end of the current service execution). This adaptation of the execution behaviour saves the cost of maintaining the search graph and triggering re-selection a potentially very large number of times. Given this, only the critical short interval $[t_{crit}, t_e]$ is relevant for the purpose of evaluating our approach, and therefore we assume next a short execution time per learning object, ignoring the irrelevant long interval $[t_s, t_{crit}]$. This simplifies the experiments and facilitates averaging the results over multiple runs. All the results reported are averaged over 30 randomly-generated requests (with two global-level constraints and a utility optimisation requirement).



Fig. 10. Reselection time in response to a violation in the executed LO's qualities.

9.2. Re-selection time results

To assess the gain in performance obtained by the proposed approach, we first study the time required for re-selecting services for the remaining tasks in response to an affecting change at execution time. We compare two strategies:

- *Reselection from scratch*. This strategy recalculates the optimal instances for the non-executed nodes from scratch, in response to a change. It is based on the multi-constrained Bellman–Ford algorithm, and is originally introduced in [11].
- *Repair-based reselection*. This is the proposed re-selection strategy in this paper, where adaptation to a change is achieved by only making the necessary updates to the optimal instances already recorded at nodes.

Here, the number of learning objects (i.e. services) considered per task is fixed at 500, while the execution position (the index of the learning object being executed) when the change occurs is varied between 1 and 10 (11 is the total number of tasks per abstract plan). Figs. 9 and 10 compare the two strategies in terms of running time, averaged over a number of different random requests. In Fig. 9, change types (addition, deletion, or changes in qualities) and locations (the tasks and services affected) at each execution position are selected randomly, whereas those considered in Fig. 10 are receiving unexpected quality values from the executed services. As can be seen, the repair-based re-selection significantly outperforms the re-selection from scratch, especially when the change is discovered at an early stage of execution. Moreover, both strategies require less time with the increasing execution position. This is because, as more services are executed, the number of remaining graph nodes to be considered in the re-selection process decreases (e.g. 10 nodes at execution position 1, 9 nodes at execution position 2, etc.), and so does the number of their optimal instances. We can also observe from the situation studied in Fig. 10, in which it is not possible to perform the adaptation process in parallel with execution since the erroneous behaviour cannot be discovered prior to its occurrence, that almost no interruption in execution will be caused by the repair-based approach proposed.

9.3. Optimality results

To assess the gain in utility obtained by the proposed approach, we compare here two strategies:

- *Delayed Reaction*. This strategy delays the reaction to a change, i.e. until when the unavailable learning object is invoked or after the quality violating learning object is executed. The re-selection of services in reaction to the change is conducted

² <http://www.openarchives.org/OAI/openarchivesprotocol.html>.




Fig. 11. Deletion of the selected learning object at position 11.




Fig. 12. Changes in the qualities of the selected learning object at position 11.




Fig. 13. Deletion of a selected learning object (Exec. Pos. = 1).

by enumerating all possible service combinations for the remaining tasks, and selecting the best possible alternative. This strategy provides an upper bound for the performance, in terms of solution optimality, of existing *reactive* approaches in the literature, including optimal re-selection reactive approaches [3], heuristic re-selection reactive approaches [5], and pre-computed backups reactive approaches [6]. In particular, the solution optimality of this strategy equals that of optimal re-selection reactive approaches, while the performance of heuristic-based and backup-based reactive re-selection ap-




Fig. 14. Changes in the qualities of a selected learning object (Exec. Pos. = 1).

proaches may be lower due to providing close-to-optimal (rather than optimal) solutions as outlined in Section 2.

- *Immediate Reaction*. This is the early reaction to changes proposed in this paper, i.e. reacting to a change in the selected learning objects as soon as it occurs in the environment.

The solution optimality is estimated as $\frac{cu_{act}}{cu_{opt}}$, where cu_{act} is the actual utility achieved by re-selecting services for the non-executed tasks, and cu_{opt} is the optimal utility assuming no task is executed (i.e. the utility of the best solution according to the current environment state, and given that no task is executed). Figs. 11 and 12 show the results in the cases where the last service in the selected solution becomes unavailable, and changes its qualities, respectively, varying the execution position at which the change occurs (i.e. at which the re-selection is triggered by the early approach to change handling) between 1 and 10 (each solution composite service is comprised of 11 services). As expected, the earlier change adaptation is performed, the better the utility of the resulting solution, which emphasises the importance of responding to changes as early as possible. This is further highlighted in Figs. 13 and 14, where the execution position at which the change is observed is fixed at 1 (i.e. the change occurs while executing the first service of the selected solution), while the change location (the index of the task affected by the change) ranges between 2 and 11. Clearly, the optimality achieved by delayed re-selection decreases as more services are executed.

9.4. Interruption time results

Finally, the last part of the experiments evaluates the reduction in interruption time between component service executions, achieved as a result of reacting to changes in the environment as soon as they occur, in parallel with the execution of the current component service. This parallelism is simulated using multi-threading on the composite service provider side, with the execution of a component service being simulated by invoking a service on a remote computer, which simply sleeps for a certain amount of time $execTime$ (service execution time) before returning a result. Changes occurring during execution are generated randomly in the interval $[start, end = start + (execTime * maxNum)]$, where $start$ is the start time of the composite service execution, while $maxNum$ is the maximum number of component services (learning objects) in a composite solution. The type of each change (addition, deletion, or changes in qualities), and its location (the task and the learning object affected by the change) are also selected randomly. Fig. 15 shows the delay time after completing the execution of each learning object in the composite solution, averaged over a number of different runs. In each run, 20




Fig. 15. Interruption time with respect to execution position.

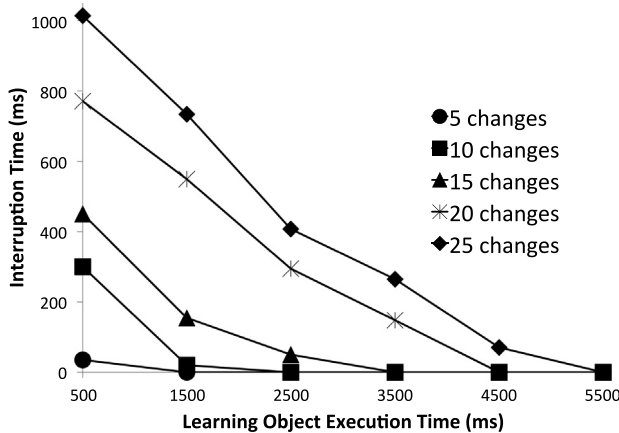


Fig. 16. The effect of the number of changes per learning object execution.

changes to be considered are introduced during each component service execution, while the number of services per task, service execution time *execTime*, and the number of tasks per abstract plan *maxNum*, are fixed at 500, 5 s, and 11 tasks, respectively. The results indicate that, even with this large number of changes, the interruption time achieved is small (i.e. the re-selection process in response to the changes is almost completed before the current component service finishes its execution), especially with the increasing execution position (the re-selection process requires less time when more component services become executed, due to the decreased number of nodes to be considered in the re-selection process).

The interruption time is further evaluated in Figs. 16 and 17 with respect to service execution time *execTime*. Fig. 16 shows the interruption time between the first and second component service executions, varying the number of changes introduced during the first service execution between 5 and 25 (all the changes are assumed to be interrupting). As expected, the interruption time decreases with the decreasing number of changes and the increasing service execution time. Fig. 17, on the other hand, shows the effect of varying the time slot within which the change occurs, on interruption time. For this purpose, the service execution time *execTime* is divided into 25 equal time intervals, during which an interrupting change is introduced while executing the first component service. Intuitively (as shown in Fig. 17), the earlier the change occurs during a service execution, the more likely that no interruption will be caused by the corresponding re-selection.

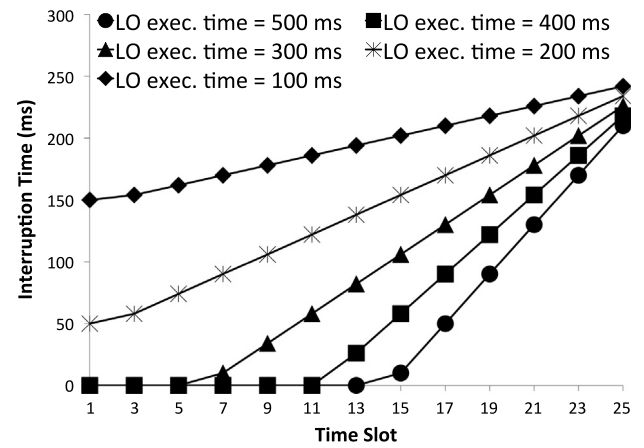


Fig. 17. The effect of the time slot within which the change occurs.

10. Conclusion and future work

In this paper, we have presented a novel adaptive execution algorithm, capable of handling execution-time service changes for both repairing and optimisation purposes. To achieve a light adaptation process, the algorithm reuses the optimal instances generated during the selection process. For this purpose, it assumes a reverse version of the search graph, which allows response to changes by applying only a minimal number of modifications to the graph, without the need to perform re-selection from scratch. The adaptation process is triggered as soon as changes occur in the environment, without interfering with the execution process, unless necessary. This need is identified based on a categorisation of changes, specifying their urgency and importance, and guiding the behaviour of the executing system. Via such an early, parallel-to-execution, and light reaction approach, the chances of a successful recovery are maximised and solution optimality is increased, while reducing execution disruptions as much as possible, as demonstrated through the evaluation conducted. In particular, the results show that, even with changes as frequent as 20 changes (to be considered) during a service execution, the interruption time remains marginal, and decreases with increasing service execution time. Moreover, in the cases where interference with execution is non-preventable (e.g. when an executed service delivers unanticipated quality values), the algorithm manages to recover from the situation with minimal interruption.

Despite generally proving to achieve efficiency, the pruning in this paper is largely dependent on the quality distributions of services and the imposed request, and hence may not scale well in some specific situations (e.g., particular quality settings where no service dominates another due to negatively related attributes). Therefore, it would be valuable to investigate suitable heuristic-based alternatives for such situations. Also, the current version of this work does not provide any support for reasoning about the degree of reliability of service offerings when performing service selection. That is, all services are assumed to have the same credibility regarding their promised quality values. We plan to investigate how trust and reputation models could be adopted in the context of our work in order to improve the service re-selection process. Moreover, we have concentrated on a forward recovery mechanism, i.e. only the tasks that are not yet executed can be re-assigned to other services or replaced with other tasks from an alternative plan. We plan to extend our recovery mechanism to allow for the possibility of rolling back the composite service execution to a previous point in time, applicable in the cases where the effects of executed services could be undone.

References

- [1] G. Canfora, M. Penta, R. Esposito, M. Villani, QoS-Aware replanning of composite web services, in: Proc. IEEE Int. Conf. on Web Services, 2005, pp. 121–129. <http://dx.doi.org/10.1109/ICWS.2005.96>.
- [2] T. Yu, K. Lin, Adaptive algorithms for finding replacement services in autonomic distributed business processes, in: Proc. Int. Symp. on Autonomous Decentralized Systems, 2005, pp. 427–434.
- [3] D. Ardagna, B. Pernici, Adaptive service composition in flexible processes, *IEEE Trans. Softw. Eng.* 33 (2007) 369–384.
- [4] R. Berbner, M. Spahn, N. Repp, O. Heckmann, R. Steinmetz, Dynamic replanning of web service workflows, in: Proc. IEEE Int. Conf. on Digital Ecosystems and Technologies, 2007, pp. 211–216.
- [5] K.J. Lin, J. Zhang, Y. Zhai, B. Xu, The design and implementation of service process reconfiguration with end-to-end QoS constraints in SOA, *Serv. Oriented Comput. Appl.* 4 (3) (2010) 157–168.
- [6] F. Wagner, B. Kloepper, F. Ishikawa, S. Honiden, Towards robust service compositions in the context of functionally diverse services, in: Proc. Int. Conf. on World Wide Web, 2012, pp. 969–978.
- [7] J. Hieslacher, R. Kazhamiakin, A. Metzger, M. Pistore, A framework for proactive self-adaptation of service-based applications based on online testing, in: Proc. Euro. Conf. on Towards a Service-Based Internet, 2008, pp. 122–133.
- [8] Y. Dai, L. Yang, B. Zhang, QoS-driven self-healing web service composition based on performance prediction, *Comput. Sci. Technol.* 24 (2009) 250–261.
- [9] R. Aschoff, A. Zisman, QoS-driven proactive adaptation of service composition, in: Proc. Int. Conf. on Service-Oriented Computing, 2011, pp. 421–435.
- [10] L. Barakat, S. Miles, M. Luck, Reactive service selection in dynamic service environments, in: Proc. Euro. Conf. on Service-Oriented and Cloud Computing, 2012, pp. 17–31.
- [11] L. Barakat, S. Miles, I. Poernomo, M. Luck, Efficient multi-granularity service composition, in: Proc. IEEE Int. Conf. on Web Services, 2011, pp. 227–234.
- [12] T. Yu, Y. Zhang, K. Lin, Efficient algorithms for web services selection with end-to-end QoS constraints, *ACM Trans. Web* 1 (1) (2007).
- [13] L. Li, J. Wei, T. Huang, High performance approach for Multi-QoS constrained web services selection, in: Proc. Int. Conf. on Service-Oriented Computing, 2007, pp. 283–294.
- [14] G. Canfora, M. Penta, R. Esposito, M. Villani, An approach for QoS-aware service composition based on genetic algorithms, in: Proc. Genetic and Evolutionary Computation Conf., 2005, pp. 1069–1075.
- [15] S. Stein, T. Payne, N. Jennings, Flexible provisioning of web service workflows, *ACM Trans. Internet Technol.* 9 (1) (2009) 1–45.
- [16] D. Ivanovic, M. Carro, M. Hermenegildo, Towards data-aware QoS-driven adaptation for service orchestrations, in: Proc. Int. Conf. on Web Services, 2010, pp. 107–114.
- [17] M. Colombo, E. Nitto, M. Mauri, SCENE: A service composition execution environment supporting dynamic changes disciplined through rules, in: Proc. Int. Conf. on Service-Oriented Computing, 2006, pp. 191–202.
- [18] R. Tolosana-Calasanz, J.A. Banares, O.F. Rana, P. Álvarez, J. Ezpeleta, A. Hoheisel, Adaptive exception handling for scientific workflows, *Concurr. Comput.: Pract. Exper.* 22 (5) (2010) 617–642.
- [19] S. Bhiri, O. Perrin, C. Godart, Ensuring required failure atomicity of composite web services, in: Proc. Int. Conf. on World Wide Web, 2005, pp. 138–147.
- [20] J. Hadad, M. Manouvrier, M. Rukoz, TQoS: Transactional and QoS-aware selection algorithm for automatic web service composition, *IEEE Trans. Serv. Comput.* 3 (1) (2010) 73–85.
- [21] L. Zeng, B. Benatallah, A.H. Ngu, M. Dumas, J. Kalagnanam, H. Chang, QoS-aware middleware for web services composition, *IEEE Trans. Softw. Eng.* 30 (5) (2004) 311–327.
- [22] G. Chaffle, K. Dasgupta, A. Kumar, S. Mittal, B. Srivastava, Adaptation in web service composition and execution, in: Proc. IEEE Int. Conf. on Web Services, 2006, pp. 549–557.
- [23] L. Yang, Y. Dai, B. Zhang, Performance prediction based EX-QoS driven approach for adaptive service composition, *Inf. Sci. Eng.* 25 (2) (2009) 345–362.
- [24] O. Sammodi, A. Metzger, X. Franch, M. Oriol, J. Marco, K. Pohl, Usage-based online testing for proactive adaptation of service-based applications, in: Proc. IEEE 35th Annual Conf. on Computer Software and Applications, 2011, pp. 582–587.
- [25] A. Zisman, J. Dooley, G. Spanoudakis, Proactive runtime service discovery, in: Proc. IEEE Int. Conf. on Services Computing, 2008, pp. 237–245.
- [26] R.R. Aschoff, A. Zisman, Proactive adaptation of service composition, in: Proc. ICSE Workshop on Software Engineering for Adaptive and Self-Managing Systems, 2012, pp. 1–10.
- [27] R. Smith, R. Davis, Frameworks for cooperation in distributed problem solving, *IEEE Trans. Syst. Man Cybern.* 11 (1981) 61–70.
- [28] J. Cardoso, J. Miller, A. Sheth, J. Arnold, Quality of service for workflows and web service processes, *Web Semant.* 1 (2004) 281–308.
- [29] X. Yuan, X. Liu, Heuristic algorithms for multi-constrained quality of service routing, *IEEE/ACM Trans. Netw.* 10 (2) (2002) 244–256.
- [30] L. Barakat, Efficient adaptive multi-granularity service composition (Ph.D. thesis), King's College London, 2013.
- [31] R. Farrell, S. Liburd, J. Thomas, Dynamic assembly of learning objects, in: Proc. Int. Conf. on World Wide Web, 2004, pp. 162–169.



Lina Barakat received the Ph.D. degree in Computer Science from King's College London in 2013, where she now works as a research associate in the Agents and Intelligent Systems group of the Department of Informatics. Her research focuses on the issues and challenges related to QoS-aware service modelling and composition in large-scale and open service-based settings.



Simon Miles is a Reader in Computer Science and head of the Agents and Intelligent Systems group at King's College London. He received his Ph.D. in Computer Science from University of Warwick, has twice been an organising committee member for the Autonomous Agents and Multi-Agent Systems conference series, co-organised the Provenance Challenge international series, and was a member of a W3C working group which published a standard for modelling and exchanging provenance data, PROV. He is investigator on a number of research projects including DIET4Elders and Justified Assessments of Service Provider Reputation, and has over 130 publications.



Michael Luck is a Professor of Computer Science at King's College London, where he is also Dean of the Faculty of Natural and Mathematical Sciences. He received his Ph.D. from University College London in 1993, and has over 200 publications, including 12 books. He was previously Coordinator of AgentLink II, the European network for Agent-Based Computing, lead author of the AgentLink Roadmaps, and a member of the Board of the International Foundation for Autonomous Agents and Multiagent Systems (2008–2014). He is a Fellow of the British Computer Society.